

***The R\*-tree:  
An Efficient and Robust Access  
Method  
for Points and Rectangles***

***Norbert Beckmann, Hans-Peter begel  
Ralf Schneider, Bernhard Seeger***

**Presented by Snehal Thakkar**

# Contents

- Terminology
- Abstract
- Introduction
- R-tree variants
- R\*-tree
- Experiments
- Conclusions

# Terminology

- **Bounding box**
  - ◆ a.k.a. Minimum Bounding Rectangle MBR
  - ◆ smallest axis-parallel rectangle enclosing the SDT value
- **Directory rectangle**
  - ◆ geometrically the MBR of the underlying rectangles
- **Margin**
  - ◆ the sum of the lengths of the edges of a rectangle
- **PAM (Point Access Method)**
  - ◆ A data structure and associated algorithms primarily to search for points defined in multidimensional space.
- **SAM (Spatial Access Method)**
  - ◆ A data structure to search for lines, polygons, etc.

# Abstract

## ■ R-tree

- ◆ based on the heuristic optimization of the **area** of the enclosing rectangle in each inner node
- ◆ heuristic: an algorithm which usually, but not always, works or which gives nearly the right answer

## ■ R\*-tree

- ◆ incorporates a combined optimization of **area**, **margin** and **overlap** of each enclosing rectangle in the directory
- ◆ outperform existing R-tree variants
- ◆ efficiently support point and spatial data at the same time
- ◆ implementation cost is only slightly higher than other R-trees

# Introduction

- **Spatial Access Methods (SAMs)**
  - ◆ approximate a complex spatial object by axis-parallel MBR
  - ◆ a complex object is represented by a limited number of bytes
  - ◆ Although a lot of information is lost, MBR of spatial objects already preserve the most essential geometric properties
    - the location of the object
    - the extension of the object in each axis
  - ◆ most popular SAM for storing rectangles: R-tree

# R-trees

- Based on the PAM B+-tree using overlapping regions
- Optimization criterion: minimize the area of each enclosing rectangle in the inner nodes
- $M$ : max. # of entries in one node
- $m$ : min. # of entries in one node ( $2 \leq m \leq M/2$ )
- Satisfy the followings:
  - ◆ the root has at least two children unless it is a leaf
  - ◆ every non-leaf node has between  $m$  and  $M$  children unless it is the root
  - ◆ every leaf node contains between  $m$  and  $M$  entries unless it is the root
  - ◆ All leaves appear on the same level

# R-trees

- **An R-tree (R\*-tree) is completely dynamic**
  - ◆ insertions and deletions can be intermixed with queries
  - ◆ no periodic global reorganization is required
- **The structure must allow overlapping directory rectangles**
  - ◆ can not guarantee that only one search path is required for an exact match query
- **This paper shows that the overlapping-regions-technique does not imply bad average retrieval performance**

# Main Problem in R-trees

- **Known parameters of good retrieval performance affect each other in a very complex way**
  - ◆ impossible to optimize one of them without influencing others
- **Data rectangles have different sizes/shapes and the directory rectangles grow and shrink dynamically**
  - ◆ the success of methods which will optimize one parameter is very uncertain
- **A heuristic approach is applied**
  - ◆ based on many different experiments carried out in a systematic framework



# Optimization Criteria

## (O1) Minimize area covered by a directory rectangle

- ◆ minimize dead space
- ◆ improve performance since decisions which paths have to be traversed can be taken on higher levels

## (O2) Minimize overlap between directory rectangles

- ◆ decrease number of paths to be traversed

## (O3) Minimize margin of a directory rectangle

- ◆ directory rectangle will be shaped more quadratic
- ◆ quadratic objects can be packed easier
- ◆ smaller directory rectangles in the level above

## (O4) Optimize storage utilization

- ◆ height of the tree will be kept low
- ◆ reduce the query cost

# Optimization

## **(+) Minimize area**

- ◆ less covering of data space, hence overlap may be reduced

## **(+) More quadratic**

- ◆ better packing, hence easier to maintain high storage utilization

## **(−) Minimize area and overlap**

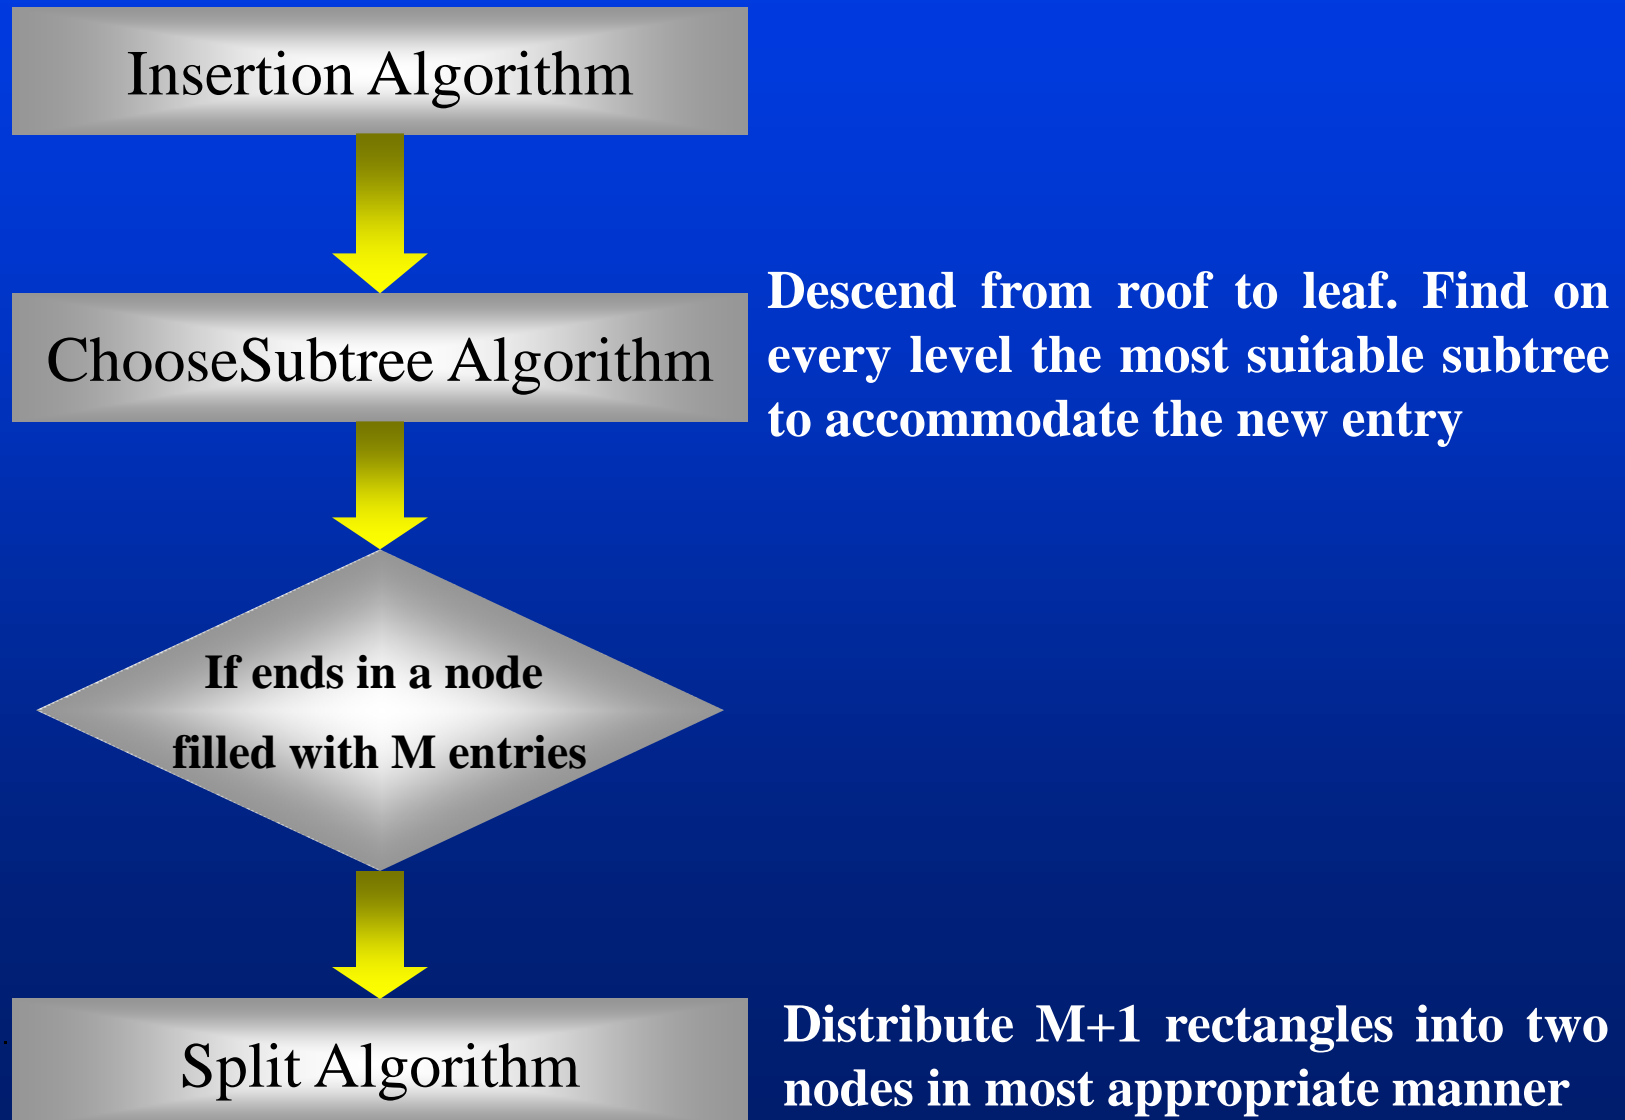
- ◆ require more freedom in # of rectangles stored in one node
- ◆ lower storage utilization
- ◆ “less quadratic” rectangles

## **(−) Minimize margin**

- ◆ reduce storage utilization

# R-tree Variants

## Optimize retrieval performance during insertion



# Original R-tree (by Guttman)

- **Method of optimization**
  - ◆ minimize area covered by a directory rectangle
  - ◆ may also reduce overlap
  - ◆ cpu cost will be relatively low

# Original R-tree (by Guttman)

## Algorithm ChooseSubtree

CS1     **[Initialize]** Set N to be the root node

CS2     **[Leaf check]**

      If N is a leaf,

          return N

      else



**[Choose subtree]**



          Choose the entry in N whose rectangle needs least area enlargement to include the new data. Resolve ties by choosing the entry with the rectangle of smallest area



      end

CS3     **[Descend until a leaf is reached]**



      Set N to be the childnode pointed to by the childpointer of the chosen entry. Repeat from CS2





















Since  $N = \text{Leaf}$ ,  
stop and return  $N$





# Split Algorithms

## Three versions

- ◆ all are designed to minimize area covered by two rectangles resulting from split

### ① Exponential

- ◆ find the area with global minimum
- ◆ CPU cost too high

### ② Quadratic and ③ Linear

- ◆ find approximation
- ◆ [Guttman] nearly same retrieval performance
- ◆ [this paper] quadratic performs much better than linear

## Algorithm QuadraticSplit

[Divide a set of  $M+1$  index entries into two groups]

QS1 [Pick first entry for each group ]

Invoke **PickSeeds** to choose two entries, each be first entry of each group

QS2 [Check if done]

Repeat

**DistributeEntry**

until

    all entries are distributed or one of the two groups has  $M-m+1$  entries (so that the other group has  $m$  entries)

QS3 [Select entry to assign ]

If entries remain, assign them to the other group so that it has the minimum number  $m$  required

## Algorithm PickSeeds

**[Choose two entries to be the first entries of the groups]**

**PS1      [Calculate inefficiency of grouping entries together]**

For each pair of entries E1 and E2, compose a rectangle R including E1 rectangle and E2 rectangle

Calculate  $d = \text{area}(R) - \text{area}(E1 \text{ rectangle}) - \text{area}(E2 \text{ rectangle})$

**PS2      [Choose the most wasteful pair ]**

Choose the pair with the largest d

**[the seeds will tend to be small, if the rectangles are of very different size (and) or the overlap between them is high]**

## Algorithm DistributeEntry

**[Assign the remaining entries by the criterion of minimum area]**



- DE1     Invoke **PickNext** to choose the next entry to be assigned
- DE2     Add It to the group whose covering rectangle will have to be enlarged least to accommodate It. Resolve ties by adding the entry to the group with the smallest area, then to the one with the fewer entries, then to either



## Algorithm PickNext



**[chooses the entry with best area-goodness-value in every situation]**



- DE1     For each entry E not yet in a group, calculate  $d_1$  = the area increase required in the covering rectangle of Group 1 to include E Rectangle. Calculate  $d_2$  analogously for Group 2
- DE2     Choose the entry with the maximum difference between  $d_1$  and  $d_2$
























N = Leaf


Node is full

*ChooseSubtree*  
*QuadraticSplit*







*ChooseSubtree*  
*QuadraticSplit*  
*PickSeeds*




E1	E2	d
J	K	
J	L	
J	M	
J	Y	
K	L	
K	M	
K	Y	
L	M	
L	Y	
M	Y	




*ChooseSubtree*  
*QuadraticSplit*  
*PickSeeds*




E1	E2	d
J	K	1.0
J	L	
J	M	
J	Y	
K	L	
K	M	
K	Y	
L	M	
L	Y	
M	Y	




*ChooseSubtree*  
*QuadraticSplit*  
*PickSeeds*



E1	E2	d
J	K	1.0
J	L	1.1
J	M	4.0
J	Y	1.7
K	L	0.1
K	M	3.6
K	Y	3.5
L	M	5.6
L	Y	4.7
M	Y	0.9





*ChooseSubtree*  
*QuadraticSplit*  
*PickSeeds*



G1	G2
L	M


E1	E2	d
J	K	1.0
J	L	1.1
J	M	4.0
J	Y	1.7
K	L	0.1
K	M	3.6
K	Y	3.5
L	Y	4.7
M	Y	0.9






*ChooseSubtree*  
*QuadraticSplit*  
*PickSeeds*  
*DistributeEntry*

G1	G2
L	M







*ChooseSubtree*  
*QuadraticSplit*  
*PickSeeds*  
*DistributeEntry*  
*PickNext*

G1	G2
L	M

Entry	d1	d2
J	1.0	
K		
Y		







*ChooseSubtree*  
*QuadraticSplit*  
*PickSeeds*  
*DistributeEntry*  
*PickNext*

G1	G2
L	M

Entry	d1	d2
J	1.0	2.0
K		
Y		







*ChooseSubtree*  
*QuadraticSplit*  
*PickSeeds*  
*DistributeEntry*  
*PickNext*

G1	G2
L	M

Entry	d1	d2
J	1.0	2.0
K	0.5	2.1
Y	2.6	0.9









*ChooseSubtree*  
*QuadraticSplit*  
*PickSeeds*  
*DistributeEntry*  
*PickNext*

G1	G2
L	M
	Y

Entry	d1	d2
J	1.0	2.0
K	0.5	2.1







*ChooseSubtree*  
*QuadraticSplit*  
*PickSeeds*  
*DistributeEntry*  
*PickNext*

G1	G2
L	M
	Y

Entry	d1	d2
J	1.0	1.2
K	0.5	2.3







*ChooseSubtree*  
*QuadraticSplit*  
*PickSeeds*  
*DistributeEntry*  
*PickNext*

G1	G2
L	M
K	Y

Entry	d1	d2
J	1.0	1.2







*ChooseSubtree*  
*QuadraticSplit*  
*PickSeeds*  
*DistributeEntry*  
*PickNext*

G1	G2
L	M
K	Y


Entry	d1	d2
J	0.7	1.2






*ChooseSubtree*  
*QuadraticSplit*  
*PickSeeds*  
*DistributeEntry*  
*PickNext*


G1	G2
L	M
K	Y
J	






*ChooseSubtree*  
*QuadraticSplit*  
*PickSeeds*  
*DistributeEntry*  
*PickNext*


G1	G2
L	M
K	Y
J	



# Problem #1: Small Seeds




# Problem #1: Small Seeds







# Problem #1: Small Seeds



X is assigned to G1.  
But the distance is  
large, and hence a bad  
split. X shall be  
assigned to G2 instead

# Problem #1: Small Seeds





# Problem #2: Prefer Bounding Rectangle



- Assign a rectangle to one seed
- Area enlarged
- Need lesser area enlargement to include next entry
- Enlarge again
- Goes on.....


G2

## Problem #2: Prefer Bounding Rectangle



- $M - m + 1 = 7$
- Uneven distribution
- Reduce storage utilization

## Problem #3



- If reached  $M-m+1$ , all remaining  $m$  entries are assigned to the other group without considering geometric properties
- Best retrieval performance when  $m = 40\%$

- $M - m + 1 = 6$
- Large overlap
- Increase number of paths to be traversed

# R-tree Variants - Greene

- **Alternative split-algorithm**
  - ◆ Use Guttman's ChooseSubtree algorithm
- **The only geometric criterion: choice of split axis**
- **More geometric optimization criteria needed to improve retrieval performance**
- **May not find “right” axis and bad split may result**

## Algorithm Greene's-Split

**[Divide a set of  $M+1$  entries into two groups]**

**GS1**     Invoke **ChooseAxis** to determine the axis perpendicular to which the split is to be performed

**GS2**     Invoke **Distribute**

# R-tree Variants - Greene


## Algorithm ChooseAxis

[Divide a set of  $M+1$  entries into two groups]

- CA1      Invoke **PickSeeds** to find two most distant rectangles of the current node
- CA2      For each axis record the separation of the two seeds
- CA3      Normalize the separations by dividing them by the length of the nodes enclosing rectangle along the appropriate axis
- CA4      Return the axis with the greatest normalized separation

## Algorithm Distribute

- D1      Sort the entries by the low value of their rectangles along the chosen axis
- D2      Assign the first  $(M+1) / 2$  entries to one group, the last  $(M+1) / 2$  entries to the other
- D3      If  $M+1$  is odd, then assign the remaining entry to the group whose enclosing rectangle will be increased least by its addition



- 56



# R\*-tree

- Previous R-trees consider only **area** parameter when choosing insertion path
- R\*-tree considers **area**, **margin** and **overlap** in different combinations

## Algorithm ChooseSubtree

CS1     Set N to be the root node

CS2     If N is a leaf,  
          return N

      else

          If the childpointers in N point to leaves

              [determine the minimum overlap cost]

              choose the entry in N whose rectangle needs least  
              overlap enlargement to include the new data rectangle.  
              Resolve ties by choosing the entry whose rectangle  
              needs least area enlargement, then the entry with the  
              rectangle of smallest area


          else

              [determine the minimum area cost]

              choose the entry in N whose rectangle needs least area  
              enlargement to include the new data rectangle. Resolve  
              ties by choosing entry with rectangle of smallest area

      end

CS3     Set N to be the childnode pointed to by the childpointer of  
          the chosen entry. Repeat from CS2



# Algorithm ChooseSubtree

- For choosing best leaf node, minimizing overlap performed slightly better (smaller overlap means smaller # of paths and hence higher performance)
- Cpu cost of determining overlap is quadratic in the number of entries
- For large node sizes, reduce calculation by modifying:  
[determine the **nearly** minimum overlap cost]

Sort the rectangles in N in increasing order of their **area** enlargement needed to include the new data rectangle.

Let A be the group of the first p entries. From the entries in A, considering all entries in N, choose the entry whose rectangle needs least **overlap** enlargement. Resolve ties as described before

# Algorithm ChooseSubtree

- For 2-D,  $p = 32$ : nearly no reduction of retrieval performance compared to unmodified
- Cpu cost remains higher than original ChooseSubtree
- Improve retrieval performance particularly in queries with small query rectangles on datafiles with non-uniformly distributed small rectangles or points
- Improve robustness


# Split of R\*-tree

- Along each axis
  - ◆ entries first sorted by lower value of their rectangles, then sorted by upper value
  - ◆ determine  $(M-2m+2)$  distributions of  $(M+1)$  entries into two groups for each sort


Distribution #	Group 1	Group 2
1	$m$	$M - m + 1$
2	$m + 1$	$M - m$
3	$m + 2$	$M - m - 1$
.....	.....	.....
$M - 2m + 2$	$M - m + 1$	$m$

e.g.  $M = 9, m = 4, 10$  entries

Distribution #	Group 1	Group 2
1	4	6
2	5	5
3	6	4



Sort #1



Sort #2

# Split of R\*-tree

- Determine goodness values for each distribution
  - ◆ three goodness values
    - area-value  $\text{area}[\text{bb}(\text{first group})] + \text{area}[\text{bb}(\text{second group})]$
    - margin-value  $\text{margin}[\text{bb}(\text{first group})] + \text{margin}[\text{bb}(\text{second group})]$
    - overlap-value  $\text{area}[\text{bb}(\text{first group}) \cap \text{bb}(\text{second group})]$
  - ◆ depend on goodness values, final distribution is determined

**bb:** bounding box of a set of rectangles



## Algorithm Split

- S1      Invoke **ChooseSplitAxis** to determine the axis, perpendicular to which the split is performed
- S2      Invoke **ChooseSplitIndex** to determine the best distribution into two groups along that axis
- S3      Distribute the entries into two groups


## Algorithm ChooseSplitAxis

- CSA1    For each axis
  - Sort the entries by the lower then by the upper value of their rectangles and determine all distributions as described above. Compute S, the sum of all **margin-values** of the different distributions
- end
- CSA2    Choose the axis with the minimum S as split axis

## Algorithm ChooseSplitIndex


- CSA1    Along the chosen split axis, choose the distribution with the minimum **overlap-value**. Resolve ties by choosing the distribution with minimum **area-value**

Distribution #1




Area-value = 1.00

Distribution #2




Distribution #3



Area-value = 1.11

# Split of R\*-tree



- Tests with  $m = 20\%$ ,  $30\%$ ,  $40\%$  and  $45\%$  of  $M$
- Best performance when  $m = 40\%$

# Another Example

Distribution #1




Figure 2a Overfilled node

Distribution #2




Figure 2a Overfilled node

Distribution #3




Figure 2a Overfilled node




Figure 2c: Split of the  $R^+$  tree  
where the splitaxis is vertical




Figure 2b- Greene's split where  
the splitaxis is horizontal




Figure 2b- Greene's split where  
the splitaxis is horizontal

# Forced Reinsert

- **R-tree and R\*-tree are nondeterministic**
  - ◆ different sequences of insertions will build up different trees
  - ◆ R-tree suffers from its old entries
- **A very local reorganization of directory rectangles is performed during a split**
  - ◆ it is rather poor
  - ◆ need a more powerful and less local instrument to reorganize the structure

# Deletion and Insert Algorithms by Guttman

## Algorithm Deletion


- D1      Invoke **Findleaf**
- D2      Remove index record E
- D3      Invoke **CondenseTree** -----> Invoke **Insert**
- D4      If necessary, make the child the new roof

## Algorithm Insert


- I1      Invoke **ChooseSubtree**
- I2      Install E, or invoke **SplitNode (QuadraticSplit)**
- I3      Invoke **AdjustTree**
- I4      If necessary, create new root




# Treating Underfilled Nodes in R-tree



# Treating Underfilled Nodes in R-tree



# Treating Underfilled Nodes in R-tree



# Forced Reinsert

- Expected that deletion and **re**insertion of old data rectangles improve retrieval performance
  - ◆ experiment showed performance improvement of 20% ~ 50% (static situation)
- To achieve dynamic reorganizations, R\*-tree forces entries to be reinserted during insertion routines
  - ◆ algorithms same as Guttman, except for **overflow** treatment

# Forced Reinsert

## Algorithm InsertData

ID1      Invoke **Insert** starting with the leaf level as a parameter, to insert a new data rectangle

## Algorithm Insert

- I1      Invoke **ChooseSubtree**, with the level as a parameter, to find an appropriate node N, in which to place the new entry E
- I2      If N has less than M entries, accommodate E in N. If N has M entries, invoke **OverflowTreatment** with the level of N as a parameter [for reinsertion or split]
- I3      If **OverflowTreatment** was called and a split was performed, propagate **OverflowTreatment** upwards if necessary  
If **OverflowTreatment** caused a split of the root, create a new root
- I4      Adjust all covering rectangles in the insertion path such that they are MBRs enclosing then children rectangles



## Algorithm OverflowTreatment

OT1     If the level is not the root level and this is the first call of  
          **OverflowTreatment** in the given level during the insertion of  
          one data rectangle, then  
              invoke **Reinsert**  
          else  
              invoke **Split**  
          end

## Algorithm Reinsert


RI1     For all  $M+1$  entries of a node  $N$ , compute the distance  
          between the centers of their rectangles and the center of the  
          bounding rectangle of  $N$   
RI2     Sort the entries in decreasing order of their distances  
          computed in RI1  
RI3     Remove the first  $p$  entries from  $N$  and adjust the bounding  
          rectangle of  $N$   
RI4     In the sort, defined in RI2, starting with the maximum  
          distance (= far reinsert) or minimum distance (= close  
          reinsert), invoke **Insert** to reinsert the entries


*ReInsert*




M = 4

p = 30% of M = 1






*ReInsert*  
*Insert*  
*ChooseSubtree*








*ReInsert*  
*Insert*  
*ChooseSubtree*

*Split is prevented!*



# Forced Reinsert

- $p = 30\%$  of  $M$  for all nodes yields best performance
- Change entries between neighboring nodes
  - ◆ decrease overlap
- Improve storage utilization
- Due to more restructuring, less splits occur
- Outer rectangles of a node are reinserted
  - ◆ shape of directory rectangles more quadratic
- Higher CPU cost (more insertion calls), alleviated by less splits

# Experimental Setup

- **Four R-tree variants**
  - ◆ R-tree with quadratic split algorithm (qua.Gut)
  - ◆ R-tree with linear split algorithm (lin.Gut)
  - ◆ Greene's variant of R-tree (Greene)
  - ◆ R\*-tree

# Experimental Setup

- Six data files containing about 100,000 2D rectangles
- Distribution of centers of rectangles
  - ◆ Uniform: 2D independent uniform distribution
  - ◆ Cluster: distribution with 640 clusters, 1600 object/cluster
  - ◆ Parcel: 100,000 disjoint rectangles. Expand each area by factor of 2.5
  - ◆ Real-Data: MBRs of elevation lines from real cartography data
  - ◆ Gaussian: 2D independent Gaussian distribution
  - ◆ Mixed-Uniform: 2D independent uniform distribution (add 1,000 large rectangles to 99,000 small rectangles, then merge)

# Experimental Setup

## ■ Types of queries

### ◆ Rectangle intersection query

- given a rectangle  $S$ , find all rectangles  $R$  in the file with  $R \cap S \neq \emptyset$

### ◆ Point query

- given a point  $P$ , find all rectangles  $R$  in the file with  $P \in R$ .

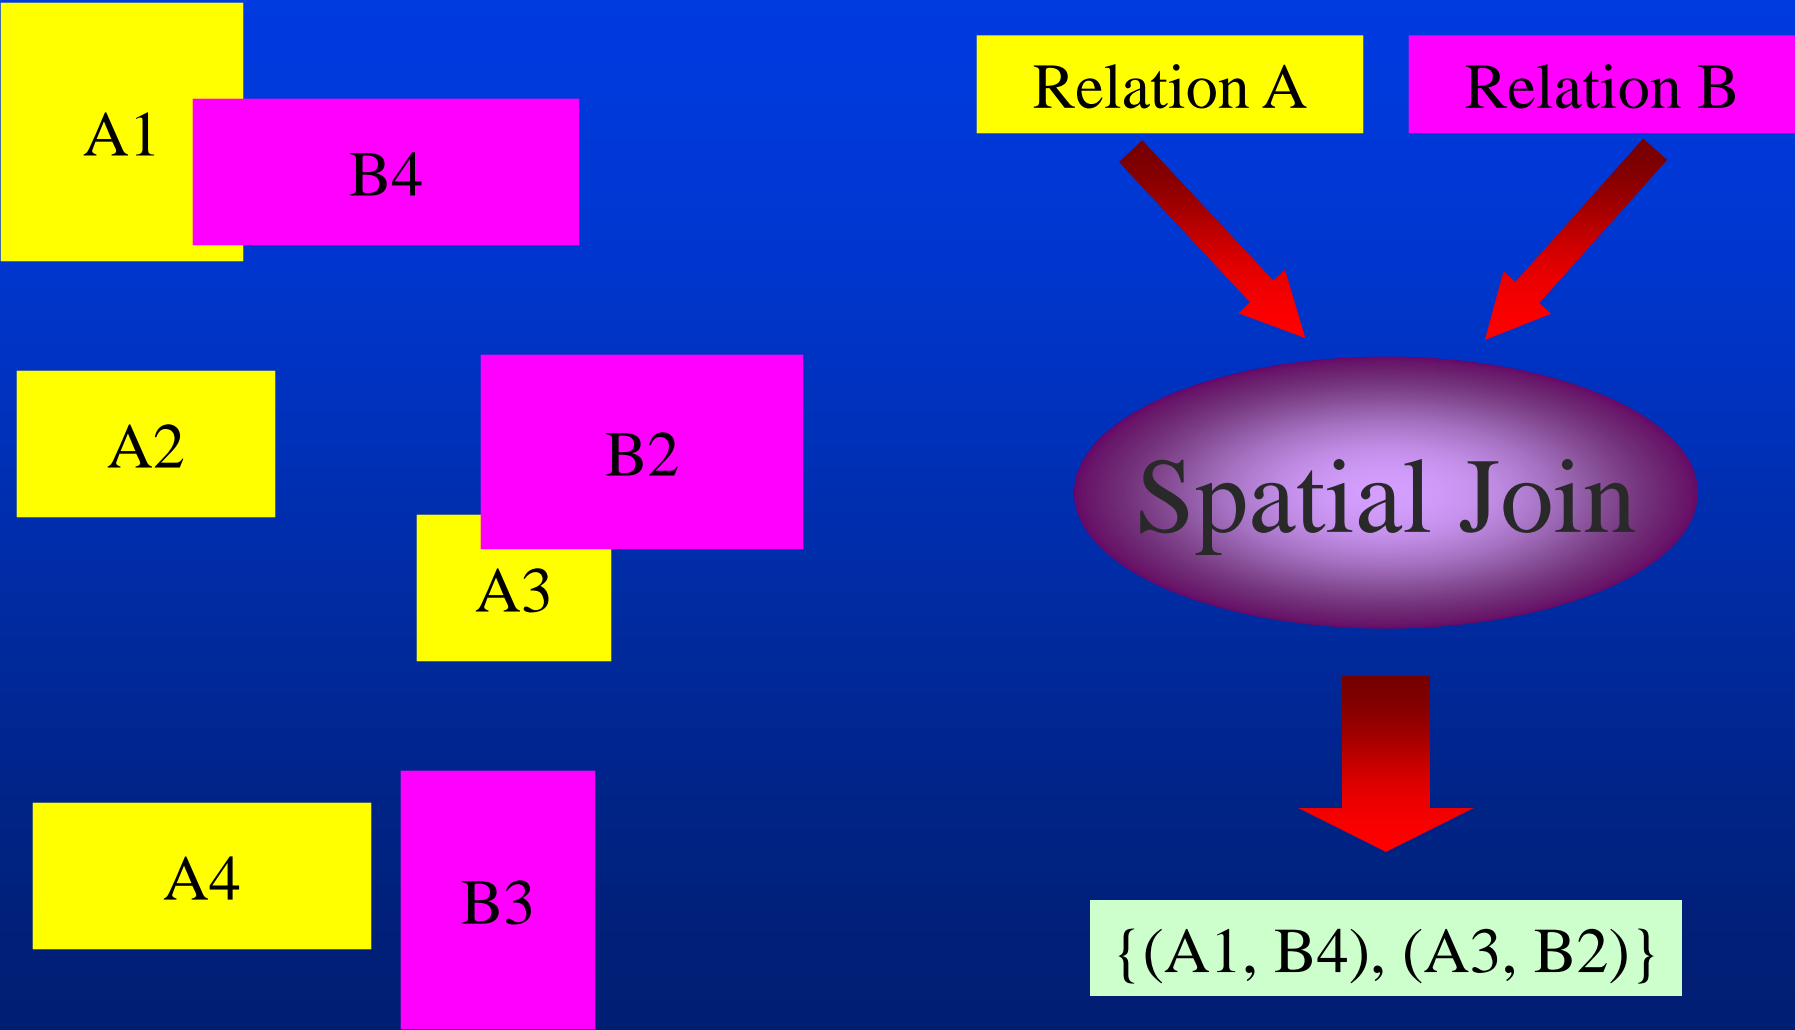
### ◆ Rectangle enclosure query

- given a rectangle  $S$ , find all rectangles  $R$  in the file with  $R \supseteq S$

### ◆ Spatial join


## ■ All experiments measured in number of disk access

# Spatial Join (Map Overlay)




# Results of Experiments

Max. performance gain of R\*-trees over all queries



Average performance gain of R\*-trees for spatial join



# Results of Experiments

- R\*-tree outperforms R-tree variants in all experiments
- Higher gain in R\*-tree for smaller query rectangles
  - ◆ storage utilization gets more important for larger rectangles
- R\*-tree has best storage utilization
- Despite use of Forced Reinsert, average insertion cost of R\*-tree is still lowest




# R\*-tree: an Efficient Point Access Method


- Another experiment on benchmark for PAM [KSSSS 89]
  - ◆ include 2-level grid file, a very popular PAM
- Performance gain considerably higher for points than for rectangles
- R\*-tree outperforms R-tree variants for point data and storage utilization, even grid file

# Results of Experiments

Average query cost averaged  
over all query and data file



Average insertion cost



# Conclusion

- **R\*-tree can efficiently be used as an access method in database systems organizing both multidimensional points and spatial data**
- **Since area, margin and overlap are reduced, R\*-tree is very robust against ugly data distribution**